

# entrenamiento\_modelos

July 11, 2025

## 1 ENTRENAMIENTO DE LOS MODELOS

```
[2]: # INSTALAR LIBRERÍAS NECESARIAS
%pip install scikit-learn
```

```
Collecting scikit-learn
  Using cached scikit_learn-1.7.0-cp312-cp312-win_amd64.whl.metadata (14 kB)
Requirement already satisfied: numpy>=1.22.0 in
c:\users\carolina\documents\proyectos_programacion\predestu\.venv\lib\site-
packages (from scikit-learn) (2.3.1)
Collecting scipy>=1.8.0 (from scikit-learn)
  Using cached scipy-1.16.0-cp312-cp312-win_amd64.whl.metadata (60 kB)
Collecting joblib>=1.2.0 (from scikit-learn)
  Using cached joblib-1.5.1-py3-none-any.whl.metadata (5.6 kB)
Collecting threadpoolctl>=3.1.0 (from scikit-learn)
  Using cached threadpoolctl-3.6.0-py3-none-any.whl.metadata (13 kB)
Using cached scikit_learn-1.7.0-cp312-cp312-win_amd64.whl (10.7 MB)
Using cached joblib-1.5.1-py3-none-any.whl (307 kB)
Using cached scipy-1.16.0-cp312-cp312-win_amd64.whl (38.4 MB)
Using cached threadpoolctl-3.6.0-py3-none-any.whl (18 kB)
Installing collected packages: threadpoolctl, scipy, joblib, scikit-learn
```

```
----- 1/4 [scipy]
----- 1/4 [scipy]
----- 1/4 [scipy]
----- 1/4 [scipy]
----- 1/4 [scipy]
----- 1/4 [scipy]
----- 1/4 [scipy]
----- 1/4 [scipy]
----- 1/4 [scipy]
----- 1/4 [scipy]
----- 1/4 [scipy]
----- 1/4 [scipy]
----- 1/4 [scipy]
----- 1/4 [scipy]
----- 1/4 [scipy]
----- 1/4 [scipy]
----- 1/4 [scipy]
```

[illegible]

[illegible]

[illegible]

[illegible]

```
Successfully installed joblib-1.5.1 scikit-learn-1.7.0 scipy-1.16.0
threadpoolctl-3.6.0
```

Note: you may need to restart the kernel to use updated packages.

```
[17]: # IMPORTACIÓN DE LIBRERÍAS
from sklearn.linear_model import LogisticRegression
from sklearn.model_selection import train_test_split
from sklearn.metrics import classification_report
import pandas as pd
import joblib
from sklearn.metrics import classification_report, confusion_matrix, accuracy_score
from sklearn.ensemble import RandomForestClassifier
```

```
[6]: # CARGAR DATASET
df_entrenamiento = pd.read_csv('../fuentes_datos/datasets_finales/CSV/
↳ dataset_entrenamiento.csv', delimiter=';')
```

## 1.1 Regresión Logística Multiclase

```
[7]: X = df_entrenamiento[['id_escuela', 'id_plan', 'id_ciclo', 'n_reprobaciones',  
    ↪ 'id_semestre', 'tutorias', 'uso_acra']]  
y = df_entrenamiento['riesgo_academico']  
  
X_train, X_test, y_train, y_test = train_test_split(X, y, test_size=0.3,  
    ↪ random_state=42)  
  
modelo_lr = LogisticRegression(multi_class='multinomial', solver='lbfgs',  
    ↪ max_iter=1000)  
modelo_lr.fit(X_train, y_train)  
y_pred_lr = modelo_lr.predict(X_test)  
  
print("=== Logistic Regression ===")  
print(classification_report(y_test, y_pred_lr))
```

```
c:\Users\carolina\Documents\Proyectos_programacion\PredEstu\.venv\Lib\site-
packages\sklearn\linear_model\_logistic.py:1264: FutureWarning: 'multi_class'
was deprecated in version 1.5 and will be removed in 1.7. From then on, it will
always use 'multinomial'. Leave it to its default value to avoid this warning.
    warnings.warn(
```

=== Logistic Regression ===					
	precision	recall	f1-score	support	
	0	0.00	0.00	0.00	40339
	1	0.00	0.00	0.00	42759
	2	0.34	0.11	0.16	82819
	3	0.33	0.90	0.48	86884
	4	0.00	0.00	0.00	13197
	accuracy			0.33	265998
	macro avg	0.13	0.20	0.13	265998
	weighted avg	0.21	0.33	0.21	265998

```
c:\Users\carolina\Documents\Proyectos_programacion\PredEstu\.venv\Lib\site-
packages\sklearn\linear_model\_logistic.py:470: ConvergenceWarning: lbfgs failed
to converge after 1000 iteration(s) (status=1):
STOP: TOTAL NO. OF ITERATIONS REACHED LIMIT
```

Increase the number of iterations to improve the convergence (max\_iter=1000).

You might also want to scale the data as shown in:

<https://scikit-learn.org/stable/modules/preprocessing.html>

Please also refer to the documentation for alternative solver options:

`https://scikit-learn.org/stable/modules/linear\_model.html#logistic-regression`

```
n iter i = check optimize result(
```

```
c:\Users\carolina\Documents\Proyectos_programacion\PredEstu\.venv\Lib\site-
packages\sklearn\metrics\_classification.py:1706: UndefinedMetricWarning:
Precision is ill-defined and being set to 0.0 in labels with no predicted
samples. Use `zero_division` parameter to control this behavior.
```

```
_warn_prf(average, modifier, f"{metric.capitalize()} is", result.shape[0])
c:\Users\carolina\Documents\Proyectos_programacion\PredEstu\.venv\Lib\site-
packages\sklearn\metrics\_classification.py:1706: UndefinedMetricWarning:
Precision is ill-defined and being set to 0.0 in labels with no predicted
samples. Use `zero_division` parameter to control this behavior.
```

```
_warn_prf(average, modifier, f"{metric.capitalize()} is", result.shape[0])
c:\Users\carolina\Documents\Proyectos_programacion\PredEstu\.venv\Lib\site-
packages\sklearn\metrics\_classification.py:1706: UndefinedMetricWarning:
Precision is ill-defined and being set to 0.0 in labels with no predicted
samples. Use `zero_division` parameter to control this behavior.
```

```
_warn_prf(average, modifier, f"{metric.capitalize()} is", result.shape[0])
```

```
[11]: # Guarda el modelo en un archivo .pkl
joblib.dump(modelo_lr, "../modelos_generados/
      ↪modelo_regresion_logistica_multiclase.pkl")
```

```
[11]: ['../modelos_generados/modelo_regresion_logistica_multiclase.pkl']
```

```
[15]: # Predicciones
y_pred = modelo_lr.predict(X_test)

# Reporte general
print("=== Evaluación del Modelo ===")
print(classification_report(y_test, y_pred))
```

```
=== Evaluación del Modelo ===
```

	precision	recall	f1-score	support
0	0.00	0.00	0.00	40339
1	0.00	0.00	0.00	42759
2	0.34	0.11	0.16	82819
3	0.33	0.90	0.48	86884
4	0.00	0.00	0.00	13197
accuracy			0.33	265998
macro avg	0.13	0.20	0.13	265998
weighted avg	0.21	0.33	0.21	265998

```
c:\Users\carolina\Documents\Proyectos_programacion\PredEstu\.venv\Lib\site-
packages\sklearn\metrics\_classification.py:1706: UndefinedMetricWarning:
Precision is ill-defined and being set to 0.0 in labels with no predicted
samples. Use `zero_division` parameter to control this behavior.
```

```
_warn_prf(average, modifier, f"{metric.capitalize()} is", result.shape[0])
c:\Users\carolina\Documents\Proyectos_programacion\PredEstu\.venv\Lib\site-
```

```

packages\sklearn\metrics\_classification.py:1706: UndefinedMetricWarning:
Precision is ill-defined and being set to 0.0 in labels with no predicted
samples. Use `zero_division` parameter to control this behavior.
    _warn_prf(average, modifier, f"{metric.capitalize()} is", result.shape[0])
c:\Users\carolina\Documents\Proyectos_programacion\PredEstu\.venv\Lib\site-
packages\sklearn\metrics\_classification.py:1706: UndefinedMetricWarning:
Precision is ill-defined and being set to 0.0 in labels with no predicted
samples. Use `zero_division` parameter to control this behavior.
    _warn_prf(average, modifier, f"{metric.capitalize()} is", result.shape[0])

```

```

[16]: # Precisión total
print("Precisión del modelo:", accuracy_score(y_test, y_pred))

```

Precisión del modelo: 0.32824306949676313

## 1.2 Random Forest Classifier

```

[18]: modelo_rf = RandomForestClassifier(n_estimators=100, random_state=42)
modelo_rf.fit(X_train, y_train)
y_pred_rf = modelo_rf.predict(X_test)

print("=== Random Forest ===")
print(classification_report(y_test, y_pred_rf))

```

```

=== Random Forest ===

```

	precision	recall	f1-score	support
0	0.50	0.39	0.44	40339
1	0.52	0.36	0.43	42759
2	0.54	0.60	0.57	82819
3	0.54	0.66	0.59	86884
4	0.57	0.30	0.39	13197
accuracy			0.53	265998
macro avg	0.53	0.46	0.48	265998
weighted avg	0.53	0.53	0.53	265998

```

[19]: # Guarda el modelo en un archivo .pkl
joblib.dump(modelo_rf, "../modelos_generados/modelo_random_forest.pkl")

```

```

[19]: ['../modelos_generados/modelo_random_forest.pkl']

```

```

[20]: # Predicciones
y_pred = modelo_rf.predict(X_test)

# Reporte general
print("=== Evaluación del Modelo ===")
print(classification_report(y_test, y_pred))

```



=== Evaluación del Modelo ===

	precision	recall	f1-score	support
0	0.50	0.39	0.44	40339
1	0.52	0.36	0.43	42759
2	0.54	0.60	0.57	82819
3	0.54	0.66	0.59	86884
4	0.57	0.30	0.39	13197
accuracy			0.53	265998
macro avg	0.53	0.46	0.48	265998
weighted avg	0.53	0.53	0.53	265998

```
[21]: # Precisión total
print("Precisión del modelo:", accuracy_score(y_test, y_pred))
```

Precisión del modelo: 0.53356792156332